

Practical Session 4 – Reinforcement Learning

A Basic Multi-Armed Bandit

We're going to start today by implementing a basic multi-armed bandit algorithm (UCB1) and exploring its behaviour in theory. We'll then see if we can tune the algorithm to converge in a good time, before applying it to our emergent software system.

Create a new project with a new file for a main component and (in a resources folder) another one for an interface UCB. In the interface definition, use the following type:

```
interface UCB {
    void setActions(String actions[])
    int getAction()
    void consumeData(dec reward)
}
```

Next we'll write our main component, which will be a simple test driver for our UCB implementation. We'll read in simulated data for a set of actions, which we can use to explore the behaviour of UCB1.

Open the source file for your main component and add the following code, which uses a JSONEncoder to read in a list of values:

```
data Row {
    dec values[]
}

component provides App requires io.File, data.json.JSONEncoder parser, UCB ucb {

    int App:main(AppParam params[])
    {
        File fd = new File("test_data.txt", File.FILE_ACCESS_READ)
        Row rows[] = parser.jsonToArray(fd.read(fd.getSize()), typeof(Row[]), null)

        String actions[]
        for (int i = 0; i < rows[0].values.arrayLength; i++)
            actions = new String[] (actions, new String("a" + i))

        ucb.setActions(actions)

        for (int i = 0; i < rows.arrayLength; i++)
        {
            int a = ucb.getAction()
            ucb.consumeData(rows[i].values[a])
        }

        return 0
    }
}
```

We now implement the UCB1 algorithm in a component `UCB.dn` as follows:

```
component provides UCB requires util.Math math {

    String learningActions[]
    dec rewards[]
    dec counts[]
    int totalCount
    int currentAction

    void UCB:setActions(String actions[]) {
        learningActions = actions
        rewards = new dec[actions.arrayLength]
        counts = new dec[actions.arrayLength]
    }

    int selectAction() {
        for (int i = 0; i < counts.arrayLength; i++) {
            if (counts[i] == 0.0) return i
        }

        dec maxVal = 0.0
        int maxInd = 0

        for (int i = 0; i < counts.arrayLength; i++) {
            // UCB1
            dec bonus = math.sqrt((2.0 * math.natlog(totalCount)) / counts[i])
            dec val = rewards[i] + bonus

            if (val > maxVal)
            {
                maxVal = val
                maxInd = i
            }
        }

        return maxInd
    }

    int UCB:getAction() {
        return currentAction
    }

    void UCB:consumeData(dec reward) {
        //update learning state
        counts[currentAction] += 1.0
        dec n = counts[currentAction]
        dec value = rewards[currentAction]
        dec newValue = ((n - 1.0) / n) * value + (1.0 / n) * reward
        rewards[currentAction] = newValue
        totalCount ++

        //choose new action
        currentAction = selectAction()
    }

}
```

Compile the entire project using a terminal window in the root directory with the command:

dnc .

Now download the [test_data.txt](#) file from the GitHub repository and run the main component.

We now need to write some code to let us see what's really happening in the learning system. In the main component, add some code to reveal how much time is spent taking each action and print this out at the end of the main function.

The results at this stage will look very skewed; this is because we're not yet *normalising* the reward data but are instead feeding the raw data directly into UCB.

UCB expects that all reward data is provided as a real number in the range [0...1] so we need to normalise our actual reward data into this range. To do this you'll need to decide on what the smallest and largest rewards are likely to be and use this information to translate the reward data into the needed range. For each reward value read in from the file, your translated value will be fed into the *consumeData* function instead of the raw reward value.

Now run your test again to see how the distribution of choices has been affected.

Experiments and Tuning

Next we explore how the UCB1 algorithm behaves in more detail. First analyse the data file to discover which of the three actions is the best one to take. Now compare this to the overall amount of time that the algorithm spends in this action.

Next, divide the execution time of the algorithm into fixed-size blocks (say, 20 actions) and extract the amount of time spent in each action per block over time.

How might we adjust the UCB1 algorithm so that it converges on the best action more quickly?

Applying Machine Learning to Emergent Software

We'll now try connecting our learning algorithm with a live emergent software system, to see if we can learn the best action to take in a real system.

In the *Practical4* folder you will find the *Learner.o* component. This component connects the UCB algorithm we just learned with a live emergent software system. For that, it connects to a component named *EmergentSys.dn*. To run the Learner, you should compile the *pal* folder with "dnc . -sp ../repository" and open three terminal windows: the first one run "dana -sp ../repository EmergentSys", the second one run "dana -sp ../repository Learner", and in the last terminal run the client. You should run each of the components in the described order and wait until the Learner is setting the Proxy in EmergentSys to start the client.

Note that for the learning algorithm to work properly you will need to convert *cost* into *reward* before sending any data into UCB. This is because, for the web server, a **lower** value of response time is better.

This is the opposite direction to that in which UCB works, where it considers a higher value to be better. How would you convert the cost value from the system to a (normalised) reward value to give to UCB?

In the Learner component source code, you will find a *costToReward* you should implement. You should consider the data you collected from running the system using Interactive Perception for your implementation.

Classifiers

With a working learning algorithm, the next thing to do is write a classifier to analyse and label each environment that is experienced by the system. For each separate environment you will need a separate instance of your learning algorithm and its associated state, so that learned actions can be stored for each environment.

How might you write a basic classifier which might be able to distinguish your different workloads? Think about how you can use the event data coming from the emergent system to create different classes of environment.

Try to implement your idea, then test it by switching between different workloads to see if different environments are detected and appropriate actions are learned for each one.

Efficient search space navigation

Another major problem with real-time machine learning is quickly navigating a large search space of possible actions to quickly converge on the best one. Our basic implementation of UCB1 does not use any approach to solve this problem.

How might you achieve this? What information could you use to make more advanced inferences about how different compositions are likely to perform if chosen? Is it possible not to have to try every possible composition and still find the best one?